

ai-harness — where things stand

all open workstreams · manager MGR-24 roster · generated 2026-09-13

FRESH

Newest source event: 2026-09-13 (manager log). main HEAD is 82f3e7b, 2026-09-09 (4 days old). Regenerate before relying on any figure — this page describes and links; it does not restate figures that live in the source docs.

Progress

Task status, worst-first. Rows **describe and link** — they do not restate figures that live in a source document (a stale number here would mislead before you reach the doc that would correct it).

10

workers tracked
(MGR-24)

3


tickets closed
(recent)

6

open workstreams

1

test-VERIFIED tasks

STATUS	TASK	PROGRESS	OWNER	LINKS
in progress	#179 report-v2 → report (this report) Sidebar multi-page review renderer; rename old report→manager-recap; Facts page; this ai-harness report is the proving case. Committed on feat/report-v2, awaiting boss-gated merge.	 90%	W10 · ai-harness-report	#179 (https://github.com/jdtogni78/ai-harness/issues/179) skills/report/ (new renderer)
in progress	#178 local semantic memory-recall vs grep Reported-done, committed on feat/memory-recall, NOT merged. Conditional adopt; 1.1GB torch footprint is the gate.	 90%	W9 · ai-harness-recall	#178 (https://github.com/jdtogni78/ai-harness/issues/178) f115aff (feat/memory-recall)
in progress	#175 deprecate usage-limit auto-resume monitor Superseded by native Claude auto-resume. Not a clean delete: OAuth/code-sessions client is shared by session_titles.py + session_fork_all.py — extract client first.	 60%	W7 · ai-harness-175	#175 (https://github.com/jdtogni78/ai-harness/issues/175) 8d83b75 (feat/175-deprecate-usage-limit)
to-do	#176 flaky test: run_marker_line() token collision procutil.py:40 token f"{time.time():.6f}. {os.getpid()}" collides within a microsecond; reproduced ~1/12 on main. Fix: monotonic counter / perf_counter_ns.	 10%	unassigned	#176 (https://github.com/jdtogni78/ai-harness/issues/176) remote_control/procutil.py:40

STATUS	TASK	PROGRESS	OWNER	LINKS
to-do	#174 new-session on-demand worktree provisioning #165 restored server spawn but --spawn worktree still refuses when the CLI env-pool is empty. Fix: provision a git worktree add, then spawn --spawn same-dir.	0%	unassigned	#174 (https://github.com/jdtogni78/ai-harness/issues/174)
done	#177 AI memory-harnesses research (POC) Reported-done as a POC (191 lines / 123 sources). Nothing committed — evidence is CLAIMED, not resolvable.	100%	W8 · ai-harness-research	#177 (https://github.com/jdtogni78/ai-harness/issues/177)

To-dos

Remaining scope per task. Struck-through items are done; open items are what's left.

in progress #179 report-v2 → report (this report)

- ☒ Build self-contained sidebar renderer (template/css/js/inliner)
- ☒ Vendor Alpine inline (offline, no CDN)
- ☒ Facts/Memory page + authoring guidance
- ☒ Rename report → manager-recap; fix [[report]] links + README
- ☒ Generate ai-harness report; screenshot + print-to-PDF verify
- ☐ Boss review + merge (gated)

in progress #178 local semantic memory-recall vs grep

- ☒ Build recall index (sqlite-vec, MiniLM) + eval harness
- ☒ Measure semantic vs grep vs hybrid over 239-file corpus
- ☐ Decide rollout (footprint gate: static-embedding A/B)
- ☐ Merge (boss-gated)

in progress #175 deprecate usage-limit auto-resume monitor

- ☐ Extract shared OAuth/code-sessions client (phase 2)
- ☐ Remove usage-limit monitor daemon (phase 3)

to-do #176 flaky test: run_marker_line() token collision

- ☐ Replace time.time() token with a monotonic counter
- ☐ Add a regression test that would have caught the collision

to-do #174 new-session on-demand worktree provisioning

- ☐ Provision own worktree when pool empty, then spawn same-dir
-
- ☐ Acceptance: empty-pool --spawn worktree provisions + spawns live

Resources

Every typed link in this report, grouped by kind. Follow the link to the source of truth.

board (1)

board [GitHub Project #2 — canonical board for ai-harness \(see Facts re: worker token scope\)](https://github.com/users/jdtogni78/projects/2) (<https://github.com/users/jdtogni78/projects/2>)

commit (2)

commit [f115aff \(feat/memory-recall\)](#) · #178 local semantic memory-recall vs grep

commit [8d83b75 \(feat/175-deprecate-usage-limit\)](#) · #175 deprecate usage-limit auto-resume monitor

file (2)

file [skills/report/ \(new renderer\)](#) · #179 report-v2 → report (this report)

file [remote_control/procutil.py:40](#) · #176 flaky test: run_marker_line() token collision

repo (1)

repo [jdtogni78/ai-harness \(public source of truth\)](https://github.com/jdtogni78/ai-harness) (<https://github.com/jdtogni78/ai-harness>)

ticket (6)

ticket [#179](#) (<https://github.com/jdtogni78/ai-harness/issues/179>) · [#179 report-v2 → report \(this report\)](#)

ticket [#178](#) (<https://github.com/jdtogni78/ai-harness/issues/178>) · [#178 local semantic memory-recall vs grep](#)

ticket [#175](#) (<https://github.com/jdtogni78/ai-harness/issues/175>) · [#175 deprecate usage-limit auto-resume monitor](#)

ticket [#177](#) (<https://github.com/jdtogni78/ai-harness/issues/177>) · [#177 AI memory-harnesses research \(POC\)](#)

ticket [#176](#) (<https://github.com/jdtogni78/ai-harness/issues/176>) · [#176 flaky test: run_marker_line\(\) token collision](#)

ticket [#174](#) (<https://github.com/jdtogni78/ai-harness/issues/174>) · [#174 new-session on-demand worktree provisioning](#)

Validation — verified vs claimed

Evidence gallery so this page stands in for the product. Each item is tagged verified (a resolvable pointer someone can open) or claimed (a bare assertion). A task with zero verified evidence renders a loud UNVERIFIED banner — never a silent pass.

in progress #179 report-v2 → report (this report)

✓ verified COMMAND render_review.py gate passed on the ai-harness model

```
$ python3 render_review.py --data ai-harness.json
Wrote ai-harness-report.html (self-contained)
[exit 0]
```

in progress #178 local semantic memory-recall vs grep

✓ verified COMMAND 5 smoke tests green (worker-reported, commit carries tests)

```
$ python -m unittest discover -s skills/memory-recall/tests
Ran 5 tests in 0.00s
[exit 0]
```

✓ verified METRIC recall@1 semantic vs grep (all 30 queries, 239-file corpus)

0.93 vs 0.57

recall@1 — recall@1 semantic vs grep (all 30 queries, 239-file corpus)

✓ verified METRIC recall@1 on HARD paraphrase queries — the #177 gap

0.89 vs 0.28

recall@1 — recall@1 on HARD paraphrase queries — the #177 gap

in progress #175 deprecate usage-limit auto-resume monitor

✓ **verified** COMMAND branch head exists on disk (worktree ai-harness-175)

```
$ git -C ~/dev/ai-harness-175 log -1 --oneline
8d83b75 remove usage-limit auto-resume monitor (#175)
[exit 0]
```

to-do #176 flaky test: run_marker_line() token collision

⚠ UNVERIFIED — no resolvable proof attached

This task's status rests on claims only. Attach a screenshot, a rendered output, a command result, or a diff before relying on it.

⚠ **claimed** COMMAND reproduction rate is a claim until a repro run is attached

```
$ (repro not re-run for this report)
~1/12 runs (claimed, from #176 body)
[exit 0]
```

to-do #174 new-session on-demand worktree provisioning

⚠ UNVERIFIED — no resolvable proof attached

This task's status rests on claims only. Attach a screenshot, a rendered output, a command result, or a diff before relying on it.

done

#177 AI memory-harnesses research (POC)

⚠ UNVERIFIED — no resolvable proof attached

This task's status rests on claims only. Attach a screenshot, a rendered output, a command result, or a diff before relying on it.

⚠ claimed

COMMAND worker-reported POC size — nothing committed, so not resolvable

```
s (no artifact on disk)
191 lines, 123 sources (claimed)
[exit 0]
```

Facts & Memory

Durable facts, conventions, and decisions that outlive a single session — the generatable version of a hand-built tracker's yellow "convention" box. This page **surfaces** memory; it is not a second source of truth. Prefer linking to the one-fact memory files under `~/ .claude/projects/*/memory/`.

When to record a fact (authoring guidance): Record a fact when a convention keeps getting undone, a decision's rationale isn't in the code, or a gotcha is rediscovered more than once. Each fact: the statement, why it exists, the date, a source link, and — where relevant — what made you record it. This page **surfaces** the one-fact memory files under `~/ .claude/projects/*/memory/`; it is not a second source of truth. Prefer linking to the memory file over restating it here.

Public jdtogni78/ai-harness is the source of truth; ai-harness-private is the archive.

Why: All work lands on the public repo; the private repo is a historical backup only.

 2026  [memory: repo_split_public_private](#)  [repo_split_public_private.md](#)

GitHub Project #2 IS the canonical board for ai-harness (45 of 60 items on the active repo; 13 are archive-repo leftovers).

Why: Worker sessions can't set a ticket's Status→In Progress because their gh token lacks the 'project' scope — a token-scope limitation, not a board/repo mismatch. Board-status writes belong to a scoped session (gh auth refresh -s project --hostname github.com, which a detached worker cannot complete interactively).

Recorded because: An earlier draft of this report MISDIAGNOSED the board as tracking the archive repo; independent verification (full 60-item listing) caught it before it reached a memory file. Verified-vs-claimed discipline working as intended.

 2026-09-13  [gh project item-list 2 --owner jdtogni78 \(60 items: 45 active / 13 archive\)](#)

The intermittent test_procutil RunMarker failure is a timestamp race, not a code regression.

Why: A microsecond-resolution token collides ~1/12 runs; re-running confirms green. Don't chase it as a real break.

Recorded because: Rediscovered more than once; tracked as #176.

 2026  [memory: ai-harness-flaky-run-marker-test](#)  [ai-harness-flaky-run-marker-test.md](#)

Concurrent lanes share one checkout — safe states are only 'in a commit' or 'in my working tree'.

Why: Gated work goes on its own branch, merges atomically as a merge-commit, with explicit-path adds — never a bare stash (shared stack).

 2026  memory: shared-repo-git-hygiene  shared-repo-git-hygiene.md

datadir ↔ secret coupling is a recorded cross-repo decision (GD-0010).

Why: Cross-cutting conventions live in DECISIONS.md (GD-NNNN); read it before re-litigating.

 2026-09  commit db644ca — docs(decisions): add GD-0010  DECISIONS.md · GD-0010

Timeline

Dispatch, close, and decision events, newest last.

- 2026-08-15
#165 merged (cb096de) · W2
new-session spawn reliability restored
- 2026-08-15
#163 merged (559c031) · W3
Telegram inbound bridge; suite 1247/1248
- 2026-09-09
#172 merged (82f3e7b) · W6
organizer/dream repo-org auditor skill
- 2026-09-09
#173 closed · -
drop --settings injection before remote-control verb
- 2026-09-10
#175 branch work (8d83b75) · W7
remove usage-limit monitor (in progress)
- 2026-09-12
#178 reported-done (f115aff) · W9
memory-recall eval; conditional adopt
- 2026-09-13
#179 dispatched + claimed · W10
report-v2 → report (this work)

Data

The complete embedded data model. Everything on every page derives from this one blob — copy or download it for extraction, or view-source the file.

```
{
  "meta": {
    "project": "ai-harness",
    "title": "ai-harness — where things stand",
    "slug": "ai-harness-report",
    "window": "all open workstreams : manager MGR-24 roster",
    "generated_at": "2026-09-13",
    "sources": [
      {
        "name": "git main",
        "detail": "git -C ~/dev/ai-harness log --oneline -40 main"
      },
      {
        "name": "worktrees",
        "detail": "git -C ~/dev/ai-harness-* branch/log"
      },
      {
        "name": "GitHub issues",
        "detail": "gh issue view <N> --repo jdtogni78/ai-harness"
      },
      {
        "name": "manager log",
        "detail": "~/ai-harness/manager/cse_01ESZBCXSSiahcjKokGeYp36.jsonl"
      },
      {
        "name": "memory index",
        "detail": "~/claude/projects/*/memory/MEMORY.md"
      }
    ],
    "freshness": {
      "level": "fresh",
      "note": "Newest source event: 2026-09-13 (manager log). main HEAD is 82f3e7b, 2026-09-13"
    }
  }
},
```